

# Integrated Assignment : April 2015

Day1: Parts 1-2

Day2: Parts 3-4

**Background:** PCA3 gene plays a role in Prostate Cancer detection due to its localized expression in prostate tissues and its over-expression in tumour tissues. This gene's expression profile makes it a useful marker that can complement the most frequently used biomarker for prostate cancer, PSA. There are cancer assays available that tests the presence of PCA3 in urine.

**Objectives:** In this assignment, we will be using a subset of the GSE22260 dataset, which consists of 30 RNA-seq tumour normal pairs, to assess the prostate cancer specific expression of the PCA3 gene.

Things to keep in mind:

- The libraries are polyA selected.
- The libraries are prepared as paired end.
- The samples are sequenced on Illumina's Genome Analyzer II.
- Each read is 36 bp long
- The average insert size is 150 bp with standard deviation of 38bp.
- We will only look at chromosome 9 in this exercise.
- Dataset is located here: <http://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE22260>
- 20 tumour and 10 normal samples are available
- For this exercise we will pick 3 matched pairs (C02,C03,C06 for tumour and N02,N03,N06 for normal). We can do more if we have time.

## PART 1 -----Obtainning Data and References -----

**Goals:**

- Obtain the files necessary for data processing
- Familiarize yourself with reference and annotation file format
- Familiarize yourself with sequence FASTQ format

```
#set your working directory
```

```
mkdir -p ~/workspace/rnaseq/integrated_assignment/  
export RNA_HOME=~/workspace/rnaseq/integrated_assignment
```

```
#copy the necessary reference and annotation files. Note, when initiating an environment variable, we don't need the $; however, everytime we call the variable, it needs to be preceeded by a $.
```

```
#make sure that the environment variable is set correctly
```

```
echo $RNA_HOME  
cp -r ~/CourseData/RNA_data/integrated_assignment_files/* $RNA_HOME  
cd $RNA_HOME
```

Q1) How many directories are there under the “refs” directory?

```
ubuntu@ip-10-182-231-187:~/workspace/rnaseq/integrated_assignment/refs$ tree
.
├── hg19
│   ├── bwt
│   │   └── 9
│   │       ├── 9.1.bt2
│   │       ├── 9.2.bt2
│   │       ├── 9.3.bt2
│   │       ├── 9.4.bt2
│   │       ├── 9.fa
│   │       ├── 9.rev.1.bt2
│   │       └── 9.rev.2.bt2
│   ├── fasta
│   │   └── 9
│   │       └── 9.fa
│   └── genes
│       └── genes_chr9.gtf
└── 6 directories, 9 files
```

Q2) How many exons does the gene PCA3 have?

```
ubuntu@ip-10-182-231-187:~/workspace/rnaseq/integrated_assignment/refs/hg19/genes$ grep PCA3 genes_chr9.gtf
9       antisense  exon       79379352      79379471      .       +       exon_id "ENSE00001600928"; exon_number "1"; gene_biotype "antisense"; gene_id "ENSG00000225937"; gene_name "PCA3"; tr
transcript_id "ENST00000412654"; transcript_name "PCA3-001"; tss_id "TSS5481";
9       antisense  exon       79397584      79397748      .       +       exon_id "ENSE00001597304"; exon_number "2"; gene_biotype "antisense"; gene_id "ENSG00000225937"; gene_name "PCA3"; tr
transcript_id "ENST00000412654"; transcript_name "PCA3-001"; tss_id "TSS5481";
9       antisense  exon       79398801      79398803      .       +       exon_id "ENSE00001693743"; exon_number "3"; gene_biotype "antisense"; gene_id "ENSG00000225937"; gene_name "PCA3"; tr
transcript_id "ENST00000412654"; transcript_name "PCA3-001"; tss_id "TSS5481";
9       antisense  exon       79399032      79402485      .       +       exon_id "ENSE00001664394"; exon_number "4"; gene_biotype "antisense"; gene_id "ENSG00000225937"; gene_name "PCA3"; tr
transcript_id "ENST00000412654"; transcript_name "PCA3-001"; tss_id "TSS5481";
```

Q3) How many cancer/normal samples do you see under the data directory?

```
ubuntu@ip-10-182-231-187:~/workspace/rnaseq/integrated_assignment/data$ tree
.
├── carcinoma_C02_read1.fasta
├── carcinoma_C02_read2.fasta
├── carcinoma_C03_read1.fasta
├── carcinoma_C03_read2.fasta
├── carcinoma_C06_read1.fasta
├── carcinoma_C06_read2.fasta
├── normal_N02_read1.fasta
├── normal_N02_read2.fasta
├── normal_N03_read1.fasta
├── normal_N03_read2.fasta
├── normal_N06_read1.fasta
├── normal_N06_read2.fasta
└── 0 directories, 12 files
```

NOTE: The fasta files you have copied above contain sequences for chr9 only. I have pre-processed those fasta files to obtain chr9 and also matched read1/read2 sequences for each of the samples. You do not need to redo this; However, I will explain below the process I went through to get them to this point.

##### FYI #####

- Access the following link: <http://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE22260>. Scroll down to select the files you want to download.

-The raw data in GEO is provided as \_map.txt. After you download the files, you can run the following command to convert them to FASTA:

```
cat GSM554076_C02_read1_map.txt | grep chr9 | cut -f1,2 | awk '{print ">"$1"\n"$2}' >
GSM554076_C02_read1_map.chr9.fasta
```

```
cat GSM554076_C02_read2_map.txt | grep chr9 | cut -f1,2 | awk '{print ">"$1"\n"$2}' >
GSM554076_C02_read2_map.chr9.fasta
```

-The second challenge was to match the reads for both read1 and read2, since the two FASTA files have different number of records.

```
for i in `cat GSM554076_C02_read2_map.chr9.fasta | grep ">";do R1=`echo ${i} | sed 's/0\2/0\1/g`;
grep -A1 $R1 GSM554076_C02_read1_map.chr9.fasta >> carcinoma_C02_read1.fasta;done;
```

```
for i in `cat carcinoma_C02._read1.fasta | grep ">";do R2=`echo ${i} | sed 's/0\1/0\2/g`; grep -A1
$R2 GSM554076_C02_read2_map.chr9.fasta >> carcinoma_C02_read2.fasta;done;
```

- Now you have two FASTA files with the same number of reads at the each end

##### FYI #####

#### **Q4) What sample has the highest number of reads?**

A) An easy way to figure out the number of reads is to make use of the command 'wc'. This command counts the number of lines in a file. Keep in mind that one sequence can be represented by multiple lines. Therefore, you need to first grep the read tag and count those.

```
>HWUSI-EAS230-R:6:58:12:550#0/1
TTTGTTTGTTTGCTTCTGTTTCCCCCAATGACTGA
```

running this command only give you 2\*readNumber  
>wc -l YourFastaFile.fasta

running this command will give you the proper readNumber  
>grep ">" YourFastaFile.fasta | wc -l

## PART 2 ----- Data alignment -----

### Goals:


- Familiarize yourself with Tophat/Bowtie alignment options
- Perform alignments
- Obtain alignment summary

### Q5) What is the value of `--mate-inner-dist`? What calculation did you do to get that answer?

A) Mate inner distance is the approximate distance between the reads. You can get this number by:

- 1) Using **insert size** estimates provided from the library preparation step. `--mate-inner-distance= insert size-2x(ReadLength)`
- 2) If you don't have that information, then you can subset the FASTA file and run a quick alignment. Plot the fragment distribution from this subset and use those numbers for the full alignment
- 3) We were told that the average **insert size** for these samples is 150 bp and the reads are 36bp long. so `--mate-inner-distance= 150-2x(36)=78=~80bp`

-remember this from our notes?



### Q6) Considering that the read length in this exercise is 36bp, what should you set the `--segment-length` to (default is 25bp)?

A) If you keep the default value of 25 bases, Tophat will split each read into 2 segments of 25bp and 11bp lengths. It is preferred to split the read into segments of equal length. Therefore, assigning `--segment-length` a value of 18 for a 36bp read is recommended. When deciding on a number, try avoiding a split that will result in a very short segment. Short segments might not be uniquely mapped and this can affect your transcript assembly process.

```
cd $RNA_HOME/  
export RNA_DATA_DIR=$RNA_HOME/data/  
echo $RNA_DATA_DIR  
mkdir -p alignments/tophat/trans_idx  
cd alignments/tophat  
export TRANS_IDX_DIR=$RNA_HOME/alignments/tophat/trans_idx/  
echo $TRANS_IDX_DIR
```

#take a minute and try to figure out what each parameter means and how we go the numbers.

```
tophat2 -p 8 --mate-inner-dist 80 --mate-std-dev 38 --segment-length 18 --rg-id=normal  
--rg-sample=normal_N02 -o normal_N02 -G $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--transcriptome-index $TRANS_IDX_DIR/ENSG_Genes $RNA_HOME/refs/hg19/bwt/9/9  
$RNA_DATA_DIR/normal_N02_read1.fasta $RNA_DATA_DIR/normal_N02_read2.fasta
```

```
tophat2 -p 8 --mate-inner-dist 80 --mate-std-dev 38 --segment-length 18 --rg-id=normal  
--rg-sample=normal_N03 -o normal_N03 -G $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--transcriptome-index $TRANS_IDX_DIR/ENSG_Genes $RNA_HOME/refs/hg19/bwt/9/9  
$RNA_DATA_DIR/normal_N03_read1.fasta $RNA_DATA_DIR/normal_N03_read2.fasta
```

```
tophat2 -p 8 --mate-inner-dist 80 --mate-std-dev 38 --segment-length 18 --rg-id=normal  
--rg-sample=normal_N06 -o normal_N06 -G $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--transcriptome-index $TRANS_IDX_DIR/ENSG_Genes $RNA_HOME/refs/hg19/bwt/9/9  
$RNA_DATA_DIR/normal_N06_read1.fasta $RNA_DATA_DIR/normal_N06_read2.fasta
```

```
tophat2 -p 8 --mate-inner-dist 80 --mate-std-dev 38 --segment-length 18 --rg-id=carcinoma  
--rg-sample=carcinoma_C02 -o carcinoma_C02 -G $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--transcriptome-index $TRANS_IDX_DIR/ENSG_Genes $RNA_HOME/refs/hg19/bwt/9/9  
$RNA_DATA_DIR/carcinoma_C02_read1.fasta $RNA_DATA_DIR/carcinoma_C02_read2.fasta
```

```
tophat2 -p 8 --mate-inner-dist 80 --mate-std-dev 38 --segment-length 18 --rg-id=carcinoma  
--rg-sample=carcinoma_C03 -o carcinoma_C03 -G $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--transcriptome-index $TRANS_IDX_DIR/ENSG_Genes $RNA_HOME/refs/hg19/bwt/9/9  
$RNA_DATA_DIR/carcinoma_C03_read1.fasta $RNA_DATA_DIR/carcinoma_C03_read2.fasta
```

```
tophat2 -p 8 --mate-inner-dist 80 --mate-std-dev 38 --segment-length 18 --rg-id=carcinoma  
--rg-sample=carcinoma_C06 -o carcinoma_C06 -G $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--transcriptome-index $TRANS_IDX_DIR/ENSG_Genes $RNA_HOME/refs/hg19/bwt/9/9  
$RNA_DATA_DIR/carcinoma_C06_read1.fasta $RNA_DATA_DIR/carcinoma_C06_read2.fasta
```

At this point, each one of your samples should have the following files:

```
ubuntu@ip-10-182-231-187:~/workspace/rnaseq/integrated_assignment/alignments/tophat/carcinoma_C02$ tree
.
├── accepted_hits.bam
├── align_summary.txt
├── deletions.bed
├── insertions.bed
├── junctions.bed
├── logs
│   ├── bam_merge_um.log
│   ├── bowtie_build.log
│   ├── bowtie_left_kept_reads.log
│   ├── bowtie_left_kept_reads.m2g_um.log
│   ├── bowtie_left_kept_reads.m2g_um_seg1.log
│   ├── bowtie_left_kept_reads.m2g_um_seg2.log
│   ├── bowtie_right_kept_reads.log
│   ├── bowtie_right_kept_reads.m2g_um.log
│   ├── bowtie_right_kept_reads.m2g_um_seg1.log
│   ├── bowtie_right_kept_reads.m2g_um_seg2.log
│   ├── gtf_juncs.log
│   ├── juncs_db.log
│   ├── long_spanning_reads.segs.log
│   ├── m2g_left_kept_reads.err
│   ├── m2g_left_kept_reads.out
│   ├── m2g_right_kept_reads.err
│   ├── m2g_right_kept_reads.out
│   ├── prep_reads.log
│   ├── reports.log
│   ├── reports.samtools_sort.log0
│   ├── run.log
│   ├── segment_juncs.log
│   └── tophat.log
├── prep_reads.info
└── unmapped.bam

1 directory, 30 files
```

### Q7) How would you obtain summary statistics for each aligned file?

A) There are many RNA-seq QC tools available that can provide you with detailed information about the quality of the aligned sample. However, for a simple summary of aligned reads counts you can use samtools flagstat:

`samtools flagstat accepted_hits.bam`

or

`samstat accepted_hits.bam`

hint: You can also look for the logs generated by Tophat. These logs provide a summary of the aligned reads.

## PART 3 ---- Expression Estimation -----

### Goals:

- Familiarize yourself with Cufflinks options
- Run Cufflinks to obtain expression values
- Obtain expression values for the gene PCA3

```
cd $RNA_HOME/  
mkdir expression  
cd expression
```

example (how to run cufflinks for one sample):

```
cufflinks -p 8 -o normal_N02 --GTF $RNA_HOME/refs/hg19/genes/genes_chr9.gtf --no-update-check  
$RNA_HOME/alignments/tophat/normal_N02/accepted_hits.bam  
cufflinks -p 8 -o normal_N03 --GTF $RNA_HOME/refs/hg19/genes/genes_chr9.gtf --no-update-check  
$RNA_HOME/alignments/tophat/normal_N03/accepted_hits.bam  
cufflinks -p 8 -o normal_N06 --GTF $RNA_HOME/refs/hg19/genes/genes_chr9.gtf --no-update-check  
$RNA_HOME/alignments/tophat/normal_N06/accepted_hits.bam
```

```
cufflinks -p 8 -o carcinoma_C02 --GTF $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--no-update-check $RNA_HOME/alignments/tophat/carcinoma_C02/accepted_hits.bam  
cufflinks -p 8 -o carcinoma_C03 --GTF $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--no-update-check $RNA_HOME/alignments/tophat/carcinoma_C03/accepted_hits.bam  
cufflinks -p 8 -o carcinoma_C06 --GTF $RNA_HOME/refs/hg19/genes/genes_chr9.gtf  
--no-update-check $RNA_HOME/alignments/tophat/carcinoma_C06/accepted_hits.bam
```

At this point, you should have the following files in your “expression” directory:

```
ubuntu@ip-10-182-231-187:~/workspace/rnaseq/integrated_assignment/expression$ tree
.
├── carcinoma_C02
│   ├── genes.fpk_tracking
│   ├── isoforms.fpk_tracking
│   ├── skipped.gtf
│   └── transcripts.gtf
├── carcinoma_C03
│   ├── genes.fpk_tracking
│   ├── isoforms.fpk_tracking
│   ├── skipped.gtf
│   └── transcripts.gtf
├── carcinoma_C06
│   ├── genes.fpk_tracking
│   ├── isoforms.fpk_tracking
│   ├── skipped.gtf
│   └── transcripts.gtf
├── normal_N02
│   ├── genes.fpk_tracking
│   ├── isoforms.fpk_tracking
│   ├── skipped.gtf
│   └── transcripts.gtf
├── normal_N03
│   ├── genes.fpk_tracking
│   ├── isoforms.fpk_tracking
│   ├── skipped.gtf
│   └── transcripts.gtf
└── normal_N06
    ├── genes.fpk_tracking
    ├── isoforms.fpk_tracking
    ├── skipped.gtf
    └── transcripts.gtf

6 directories, 24 files
```

**Q8) How do you get the expression of PCA3 across the normal and carcinoma samples?**

A) Cufflinks generates two expression files: gene level expression and isoform level expression. To look for the expression value of a specific gene, you can use the command ‘grep’ followed by the gene name and the path to the expression file

`grep PCA3 ./*/genes.fpk_tracking`

```
ubuntu@ip-10-182-231-187:~/workspace/rnaseq/integrated_assignment/expression$ grep PCA3 ./*/genes.fpk_tracking
./carcinoma_C02/genes.fpk_tracking:ENSG00000225937 - - ENSG00000225937 PCA3 TSS5481 9:79379351-79402485 - - 11.3862 1.20208 21.5704 OK
./carcinoma_C03/genes.fpk_tracking:ENSG00000225937 - - ENSG00000225937 PCA3 TSS5481 9:79379351-79402485 - - 167.061 121.064 213.058 OK
./carcinoma_C06/genes.fpk_tracking:ENSG00000225937 - - ENSG00000225937 PCA3 TSS5481 9:79379351-79402485 - - 761.04 686.414 835.666 OK
./normal_N02/genes.fpk_tracking:ENSG00000225937 - - ENSG00000225937 PCA3 TSS5481 9:79379351-79402485 - - 0 0 0 OK
./normal_N03/genes.fpk_tracking:ENSG00000225937 - - ENSG00000225937 PCA3 TSS5481 9:79379351-79402485 - - 65.6939 37.0228 94.3651 OK
./normal_N06/genes.fpk_tracking:ENSG00000225937 - - ENSG00000225937 PCA3 TSS5481 9:79379351-79402485 - - 511.402 429.774 593.031 OK
```



## PART 4 -- Differential Expression Analysis ---

### Goals:

- Perform differential analysis between tumor and normal samples
- Check if PCA3 is differentially expressed

```
cd $RNA_HOME/expression
```

```
ls -1 */transcripts.gtf > assembly_GTF_list.txt
```

```
cuffmerge -p 8 -o merged -g $RNA_HOME/refs/hg19/genes/genes_chr9.gtf -s  
$RNA_HOME/refs/hg19/bwt/9/ assembly_GTF_list.txt
```

```
cd $RNA_HOME/  
mkdir de  
mkdir de/reference_only  
cd $RNA_HOME/alignments/tophat
```

#run cuffdiff to perform comparison

```
cuffdiff -p 8 -L Normal,Carcinoma -o $RNA_HOME/de/reference_only/ --no-update-check  
$RNA_HOME/expression/merged/merged.gtf  
normal_N02/accepted_hits.bam,normal_N03/accepted_hits.bam,normal_N06/accepted_hits.bam  
carcinoma_C02/accepted_hits.bam,carcinoma_C03/accepted_hits.bam,carcinoma_C06/accepted_hits.b  
am
```

At this point, you should have the following files under your “de” directory:

```
ubuntu@ip-10-182-231-187:~/workspace/rnaseq/integrated_assignment/de$ tree  
├── reference_only  
│   ├── bias_params.info  
│   ├── cds.count_tracking  
│   ├── cds.diff  
│   ├── cds_exp.diff  
│   ├── cds.fpk_tracking  
│   ├── cds.read_group_tracking  
│   ├── gene_exp.diff  
│   ├── genes.count_tracking  
│   ├── genes.fpk_tracking  
│   ├── genes.read_group_tracking  
│   ├── isoform_exp.diff  
│   ├── isoforms.count_tracking  
│   ├── isoforms.fpk_tracking  
│   ├── isoforms.read_group_tracking  
│   ├── promoters.diff  
│   ├── read_groups.info  
│   ├── run.info  
│   ├── splicing.diff  
│   ├── tss_group_exp.diff  
│   ├── tss_groups.count_tracking  
│   ├── tss_groups.fpk_tracking  
│   ├── tss_groups.read_group_tracking  
│   └── var_model.info  
1 directory, 23 files
```

**Q9) any significant genes that are differentially expressed? what about PCA3?**

A) Due to the small sample size, the PCA3 signal is not significant at the adjusted p-value level. You can try re-running the above exercise on your own by using all of the samples in the original data set. Does including more samples change the results?

NOTE: Make a copy of the data to use in generateCummerbund plots generation

```
cd $RNA_HOME/  
mkdir final_results  
cd $RNA_HOME/final_results  
mkdir reference_only  
cp $RNA_HOME/de/reference_only/isoform* reference_only/  
cp $RNA_HOME/de/reference_only/read_groups.info reference_only/
```

NOTE: Rerun Obi's CummerBund Script focusing on PCA3 genes.

**Q10) What plots can you generate to help you visualize this gene's expression profile?**

A) The CummerBund package provides a wide variety of plots that can be used to visualize a gene's expression profile or genes that are differentially expressed. Some of these plots include heatmaps, boxplots, and volcano plots.